

HASHCASE / APPLIED AI

# AI that gets the work done.

End-to-end workflow automation and  
private AI systems built around your  
operations.



CLIENT CAPABILITIES / SEPTEMBER 2026



# From scattered inputs to approved outcomes.

HashCase designs the complete path from source material to a finished, reviewable business result.

---

## READ THE WORK

Documents, audio, images and connected systems

---

## APPLY JUDGMENT

Grounded models, rules and staff review


---

## FINISH THE TASK

Approved outputs, records and actions

# A routine request becomes five handoffs

Illustrative example: a request arrives by email with a scanned form and a voice note.



**The hard part is locating evidence, checking rules, getting approval and updating the system — all in one controlled run.**

# Four gaps behind one unfinished task

**01**

## Fragmented intake

Files and messages arrive without a reliable, shared case record.

**02**

## Weak evidence

Teams need to verify names, numbers and source claims manually.

**03**

## Disconnected action

A draft answer still needs staff to move it between systems.

**04**

## Unclear boundary

Sensitive material needs explicit rules for access and model use.

**The workflow must carry evidence, approval and action together.**

# Two capabilities. One operating standard.

01

## Applied AI workflow automation

Intake → evidence → model and rules → review  
→ approved action

02


## Private model deployment and tuning

Inference, retrieval, data, logs and lifecycle inside a  
client-controlled boundary

**Built as products your team can run, inspect and improve.**

# An automation that carries work to completion

Every stage has an input, an output, a failure path and an owner.



**Built to handle retries, invalid inputs and partial failure — with state your team can see.**

# Built around your documents, decisions and systems

|                                |                                                   |   |                                                                           |
|--------------------------------|---------------------------------------------------|---|---------------------------------------------------------------------------|
| <b>DOCUMENT OPERATIONS</b>     | PDFs, scans, forms and attachments                | → | <b>Validated fields, review-ready records and exception queues</b>        |
| <b>KNOWLEDGE WORK</b>          | Approved internal policies and reference material | → | <b>Sourced answers, reports and reusable institutional knowledge</b>      |
| <b>MULTI-SYSTEM OPERATIONS</b> | Requests, business APIs and approval chains       | → | <b>Tracked actions, recovery from partial failure and clear ownership</b> |
| <b>MULTIMODAL PRODUCTION</b>   | Audio, notes, images and source documents         | → | <b>Multilingual text and approved visual or video assets</b>              |



# Your AI stays inside your boundary.

We deploy inference, retrieval and approved model adaptation in infrastructure you control.


Data paths, identity, network egress and logs are specified and tested for the deployment.





# A data boundary you can inspect

The exact boundary is mapped to your infrastructure, access model and approved integrations.



Unapproved external model calls are blocked by policy when the agreed private boundary requires it.

# Quality is tested at every layer

|    |               |                                                               |
|----|---------------|---------------------------------------------------------------|
| 01 | PROGRAMMATIC  | Schemas, identifiers, numbers and safe state transitions      |
| 02 | CONTENT       | Source coverage, unsupported claims, terminology and language |
| 03 | TRAJECTORY    | Tool choices, arguments, loops and final system state         |
| 04 | FAILURE CASES | Conflicts, malformed inputs, timeouts and recovery            |
| 05 | OPERATIONS    | Latency, throughput, task cost and regression after change    |

The client sees acceptance evidence, release criteria and the rollback path.

# Keep consequential decisions with your team

We place approval gates where a wrong fact or outward action would matter.

**01**

## **CORRECT THE SOURCE**

Staff inspect extracted text, names and fields before generation.

**02**

## **REVIEW THE RESULT**

The draft carries evidence, exceptions and editable output.

**03**


## **APPROVE THE ACTION**

Publication or system changes follow the client's authority rules.

**The approval design can differ by task: staff may review every output, only exceptions, or only outward actions.**

# Official inputs to reviewed public communication

For Maharashtra Government DGIPR, HashCase built a Marathi-first platform in production use.



**Safeguards include tracked jobs, verified names and terminology, source checks, schema rules and officer approval.**

Current DGIPR workflows use approved external AI services for several tasks. Private deployment is a separate offering.

# Start with one high-value workflow

We agree on the work to automate, the data boundary and the evidence required before expansion.

**01**

## MAP

Source systems, decisions, handoffs and cost of errors

---

Workflow map + acceptance cases

**02**

## BUILD

Intake, model, rules, review and integration in a working pilot

---

Usable interface + connected pipeline

**03**

## PROVE

Representative and edge cases against the current baseline

---

Evaluation record + failure modes

**04**

## OPERATE

Logging, versioning, rollback and ownership

---

Release plan + handover

**The outcome is a working process and a documented decision on how to scale it.**

WHAT WE CAN BUILD TOGETHER

# Bring us one workflow worth fixing.

We will map it, build a working pilot and  
prove it against your real cases.



taaha@hashcase.tech · [www.hashcase.tech](http://www.hashcase.tech)

