

CONFIDENTIAL - CLIENT REVIEW DRAFT

High-Value Customer Identification & Social Influence Mapping

An evidence-led account of the reported workflow for resolving customer identities, assessing public-service exposure, and routing priority support across a large brokerage customer base.

REPORTED ENGAGEMENT

Q3-Q4 2024

Pilot run and production baseline, as stated in the source case study.

DOCUMENT STATUS

Evidence-led edition v3.0

Prepared September 2026. External release requires the validation items in the appendix.

EXECUTIVE VIEW

A decision-support layer for service prioritisation

The reported engagement addressed a specific operational risk: a customer service failure involving a publicly influential customer can create exposure far beyond the value of the individual account. The proposed response combined identity resolution, influence scoring, human review, and CRM routing.



FROM OPERATIONAL PROBLEM TO DELIVERY RESPONSE

Problem observed

- KYC data is an onboarding snapshot, not a current view of public authority or reach.
- Common Indian names create high identity-collision risk.
- Raw follower count is not a reliable proxy for reputational exposure.
- Frontline teams need an executable queue, not another dashboard score.

Reported delivery response

- A staged funnel to reduce expensive public-data discovery.
- Multi-attribute corroboration before accepting a match.
- A composite score balancing reach, authority, engagement, and relevance.
- Human review for high-impact ambiguous cases.
- Tier-based support routing with override and audit controls.


EVIDENCE POSTURE

The supplied case study supports the methodology and intended operating model. It does not contain deployment records, model-validation results, achieved SLA data, client sign-off, or quantified business outcomes. This edition keeps those boundaries explicit.

DELIVERY CHRONOLOGY

From discovery to an operational baseline

The sequence below reconstructs the engagement logic from the source narrative. Dates and completion status should be matched to project records before external publication.



DELIVERY ARTEFACTS DESCRIBED

DATA

Normalized customer anchors

Names, addresses, email domains, occupation, income/vintage and relationship flags.

MODEL

Resolution and scoring logic

Weighted corroboration, confidence thresholds, and the Composite Influence Score.

OPS

Operational tiering

CRM labels, queue rules, escalation paths, target SLAs, and supervisor override.

GOV

Human and audit controls

HITL review for ambiguous matches, weekly QA sampling, and incident escalation.


Chronology evidence required

Statement of work or purchase order, kickoff minutes, sprint or release records, pilot acceptance, production change record, and the name/role of the client sponsor who approved progression between gates.

SOLUTION ARCHITECTURE

A precision funnel - not a full-base social crawl

The architecture narrows the population before expensive discovery, separates identity confidence from influence, and sends only actionable outputs to frontline systems.



DATA BOUNDARY

PERMITTED INPUT Only sources with a recorded business purpose, access right, and retention rule.	CONTROLLED MATCH No name-only acceptance. Ambiguous high-impact candidates enter human review.	MINIMUM OUTPUT CRM receives the tier, confidence, rationale, and review status - not an uncontrolled data dump.
--	--	---

External release note: replace generic source labels with the exact licensed or authorized sources actually used.

IDENTITY RESOLUTION

Confidence comes from corroboration, not name similarity

The source model combines five corroboration vectors. A false positive is treated as more damaging than a missed match, so confidence thresholds govern automation and review.

CORROBORATION VECTOR	WEIGHT	VALIDATION LOGIC
Employer / organisation consistency	35%	Fuzzy match against employer history or verified directorships
Username / email-handle overlap	20%	String similarity between customer email handle and public identifier
Geographic / city alignment	20%	Exact or metro-cluster consistency
Chronological consistency	15%	Plausibility check across age, education, and career milestones
Name rarity weight	10%	Inverse frequency across the customer population

DECISION THRESHOLDS

>= 0.85

Automated acceptance

High-confidence match; eligible for scoring.

0.70-0.84

Human review

Review only where public reach is operationally material.

< 0.70

Automated pruning

Discarded; no downstream influence profile.

MODEL VALIDATION REQUIRED BEFORE EXTERNAL RELEASE

Provide the labelled audit sample, precision by threshold, false-match rate, review volume, reviewer agreement, and error analysis for common names. None of these measurements appears in the supplied source PDF.

INFLUENCE SCORING

A transparent score with an explicit relevance modifier

Only confirmed identities proceed to influence scoring. The score is intended to support service-risk routing, not marketing eligibility, credit decisions, or product suitability.

COMPOSITE INFLUENCE SCORE

$$\text{CIS} = \text{MIN}(100, [0.40 \text{ Reach} + 0.35 \text{ Authority} + 0.25 \text{ Engagement}] \times \text{Relevance})$$

Reported formula. Calibration and production validation were not supplied.

**Reach**

Log-compressed audience scale limits outlier dominance.

**Authority**

Credential hierarchy, institutional role, and third-party verification.

**Engagement**

Recent median interaction rate rather than passive audience size.

**Relevance**

Topical modifier for finance, business, or unrelated audiences.

GOVERNANCE CONTROLS**Controls that make the score operational**

- Store component scores and source timestamps, not only the final CIS.
- Expire stale public signals and re-verify material changes.
- Keep identity confidence separate from influence strength.
- Require human override reasons and retain an audit trail.
- Test segment-level error rates and disparate operational impact.

Intended and excluded uses**INTENDED**

Priority awareness, service escalation, incident communication, and post-case quality review.

EXCLUDED

Sales targeting, credit or margin decisions, account restrictions, employee screening, or unsupported surveillance.

OPERATIONAL ARCHETYPES

Five patterns used to anticipate escalation channels

Archetypes describe service-response patterns. They should not be treated as immutable customer labels or used to infer protected characteristics.

01

Startup founders & tech CXOs

PRIMARY ESCALATION CHANNEL

LinkedIn and X

SERVICE-RISK VECTOR

Fast peer-network amplification

02

Finfluencers & trading educators

PRIMARY ESCALATION CHANNEL

YouTube, Instagram, X

SERVICE-RISK VECTOR

Immediate retail-audience exposure

03

Financial journalists & editors

PRIMARY ESCALATION CHANNEL

Editorial and professional networks

SERVICE-RISK VECTOR

Mainstream amplification and inquiry

04

Senior legal & regulatory figures

PRIMARY ESCALATION CHANNEL

Formal notices and institutional escalation

SERVICE-RISK VECTOR

Compliance and governance scrutiny

05

Derivatives community leaders

PRIMARY ESCALATION CHANNEL

Private groups and community broadcasts

SERVICE-RISK VECTOR

Coordinated support-volume pressure

RESPONSIBLE-USE NOTE

Archetype assignment should support faster service, never reduced service for customers outside a priority segment.

OPERATIONAL TIERING

Translate scores into clear support actions

Continuous scores are difficult for frontline teams to execute. The reported design maps them to four service tiers with explicit routing, target response times, and escalation rules.

TIER	COHORT	ROUTING	RESPONSE	RESOLUTION
TIER 1	Executive VIP	Senior concierge desk; Head of CX alert; VP closure review	< 15 min	< 4 hr
TIER 2	High priority	Level-2 senior support; daily supervisor reconciliation	< 1 hr	< 12 hr
TIER 3	Monitored	Standard queue with trigger-based elevation	< 4 hr	24 hr
STANDARD	General retail	Standard IVR/chatbot triage and round-robin support	Standard	24-48 hr

OPERATIONAL SAFEGUARDS

OVERRIDE

Relationship managers can appeal a tier with a recorded justification.

INCIDENT

During a market-hours outage, priority cohorts receive proactive status updates.

AUDIT

Weekly QA sampling checks SLA adherence, rationale quality, and escalation efficacy.

All times shown are target SLAs from the source document. Achieved performance data was not supplied.

Design controls must be evidenced, not asserted

The source document used a blanket compliance statement. This edition replaces that assertion with a control framework that can be verified by the client legal, privacy, information-security, and compliance teams.

01

Purpose limitation

Document the service-risk purpose and prohibit unrelated reuse.

02

Source authorization

Record contract, API programme, consent, or other access basis for every external source.

03

Data minimization

Use the smallest corroboration set required; avoid uncontrolled profile replication.

04

Security boundary

Pseudonymized processing, encryption, RBAC, access logs, and environment separation.

05

Retention and deletion

Set expiry by data class; remove rejected candidates and stale enrichment.

06

Human accountability

HITL review, supervisor override, audit sampling, and an appeal path.

DO NOT CLAIM WITHOUT DOCUMENTARY SUPPORT

- Mass LinkedIn profile retrieval or storage unless the exact approved programme and rights are documented.
- DPDP or SEBI compliance without a legal/compliance review record tied to the actual processing design.
- Production deployment without change records, acceptance evidence, or operational logs.

INTERNAL VALIDATION

Evidence scorecard for every material claim

This appendix is an editorial control, not client-facing proof. Replace each open item with a redacted exhibit or remove the related claim before publication.

MATERIAL CLAIM	MINIMUM SUPPORT	STATUS
Engagement occurred	SOW / PO / kickoff record	OPEN
50-lakh population	Dataset manifest / processing log	OPEN
Pilot completed	Pilot report / acceptance email	OPEN
Production baseline	Release / change / go-live record	OPEN
CRM integration	Schema, screenshot, API or workflow evidence	OPEN
Match quality	Labelled audit sample and validation report	OPEN
Operational impact	Before/after KPI extract	OPEN
Compliance	Privacy, legal, security and compliance sign-off	OPEN
Client endorsement	Approved quotation or reference consent	OPEN

RELEASE RULE

A claim may be marked verified only when an attributable artefact exists and the project owner confirms that the redacted wording remains accurate.

INTERNAL VALIDATION

Build a compact due-diligence pack

A polished case study creates interest. A compact evidence pack closes doubt. The pack can remain confidential and be shared only under NDA or during reference checks.

A**Commercial proof**

Redacted SOW, PO, invoice, or engagement letter showing parties and dates.

B**Delivery chronology**

Kickoff, milestone, pilot acceptance, go-live, and handover records.

C**Technical artefact**

Architecture, data-flow, schema, code/release reference, or API contract excerpt.

D**Product evidence**

Redacted CRM screen, reviewer queue, dashboard, or sample decision record.

E**Validation evidence**

Sampling design, confusion matrix, false-positive analysis, and reviewer agreement.

F**Operating evidence**

Ticket routing log, SLA report, incident workflow, or QA audit sample.

G**Governance evidence**

DPIA/privacy review, source authorization, infosec review, retention schedule.

H**Client corroboration**

Approved quote, named sponsor, or reference-call permission.

SAFE REDACTION

- Remove customer names, account identifiers, credentials, URLs, and raw personal data.
- Keep dates, volumes, role names, system names, and acceptance language where permitted.
- Hash or mask record identifiers consistently so reviewers can see continuity across exhibits.
- Add an exhibit index and identify who holds the unredacted original.

RELEASE GATE

Publish only the strongest supportable version

The final external edition should be shorter than this review draft. Retain the case-study narrative and verified outcomes; remove the internal evidence pages after every material claim has been resolved.

- ☐ **01** Client name and logo use are authorized.
- ☐ **02** Engagement dates agree with source records and PDF metadata is intentional.
- ☐ **03** Every named data source was actually used and contractually permitted.
- ☐ **04** Pilot and production language is supported by acceptance or release evidence.
- ☐ **05** Targets are not presented as achieved outcomes.
- ☐ **06** Model-quality and business-impact figures include population and measurement period.
- ☐ **07** Privacy, security, legal, and compliance wording has owner approval.
- ☐ **08** Screenshots are real, current, redacted, and legible.
- ☐ **09** The client quote is approved verbatim with name and role.
- ☐ **10** A reference-check route exists for serious prospects.

DOCUMENT CONTROL**Version 3.0 | Evidence-led client-review draft**

Prepared September 2026 from the supplied five-page project narrative.

Status: not approved for external distribution until the release gate is complete.