

High-Value Customer Identification & Social Influence Mapping

A data-led approach to identifying influential customers and enabling faster escalation when service issues carry wider reputational risk.

CLIENT ORGANIZATION	DELIVERY PARTNER
Axis Securities Limited (AxisDirect)	HashCase (HashByte Pvt Ltd)
ENGAGEMENT SCOPE	DOCUMENT STATUS
Customer identification, influence assessment, and priority service routing	Confidential - client review draft
CUSTOMER BASE	PRIMARY USERS
Approximately 50 lakh KYC records	Customer service, relationship management, and operations teams

Executive summary

AxisDirect serves customers across retail equity, derivatives, currency, mutual funds, and advisory products. Some customers also have a significant public presence: they may be company founders, finance creators, journalists, senior professionals, or leaders of trading communities. When one of these customers faces an unresolved service problem, the issue can quickly attract attention beyond the individual account.

The project created a structured way to identify these customers without relying on follower count alone. We first narrowed the customer base using information already available in KYC and CRM records. We then checked each potential match against several independent details, including employer, location, email handle, and career history. Only profiles that passed the identity check were assessed for public reach and influence.

The final output is a service-priority record for customer support teams. It contains the verified customer match, the reasons behind the assessment, the assigned service tier, and the escalation path to follow. It is designed for customer service and incident response, not for marketing, credit, or sales decisions.

Engagement profile

Scale	Solution	Operational use
Approximately 50 lakh customer records	A five-stage workflow that narrows the population, verifies identity, measures influence, and assigns a service tier	Priority routing, faster escalation, proactive incident updates, and supervisor review
Five customer archetypes	Identity checks based on employer, handle, location, career history, and name rarity	Four service tiers with clear response and resolution targets
Human review for uncertain matches	A separate influence score based on reach, authority, engagement, and relevance	A recorded rationale, expiry date, and override history in the CRM workflow

What changed for the service team

Before this workflow, a high-impact customer could enter the same queue as any other support case, and the wider implications of the issue might only become clear after a public complaint appeared. The new workflow adds that context earlier. It gives the service team a consistent way to recognize the customer, understand why the case may require closer attention, and route it to the right level of support.

Business problem and project scope

A service issue involving a well-connected founder, finance creator, journalist, legal professional, or trading-community leader can spread quickly across professional and social networks. A failed order, margin dispute, delayed KYC unlock, or unresolved complaint may be seen by thousands of investors before the normal support process has reached a resolution.

The identity challenge

The first challenge was determining whether a public profile actually belonged to the AxisDirect customer. Many Indian names occur hundreds or thousands of times, and name similarity by itself is not reliable. KYC information can also be old: the customer may have changed employer, city, title, or public role since opening the account. We therefore treated the KYC record as a starting point and required several details to agree before confirming a match.

The operational challenge

A numerical score is useful only if it leads to a clear action. The project therefore converted verified customer assessments into four service tiers. Each tier defines who receives the case, how quickly the team should respond, when the case must be escalated, and who is responsible for reviewing closure.

Project objectives

- Reduce the customer population before making expensive external-data requests.
- Confirm identity using several independent details rather than a name-only match.
- Measure influence using audience size, professional standing, engagement, and relevance to financial markets.
- Convert the assessment into specific CRM actions for service and incident management.
- Keep uncertain matches, overrides, and important decisions subject to human review.

Project phases

Phase	Workstream	Work completed
1	Discovery and data review	Defined the use case, reviewed available KYC and CRM fields, mapped external-data requirements, and agreed how the output would be used by service teams.
2	Data and model setup	Standardized names and locations, created the candidate shortlist, defined identity checks, and designed the confidence and influence calculations.
3	Pilot validation	Tested candidate discovery, reviewed uncertain matches, checked the score outputs, and refined the operational tiers.
4	Operational rollout	Connected the tier output to the service workflow, defined escalation rules, and established override and quality-review procedures.

How the customer identification workflow works

Processing every customer against every external source would be costly and would create a large number of weak matches. We therefore used a staged workflow. Each stage reduces the population or adds new evidence, so the most detailed checks are applied only where they are needed.

Stage	What happens	Why it matters
1	Names are standardized across spelling and transliteration variants. Addresses are split into city, state, and PIN-code components. Corporate email domains are separated from free webmail, and customer identifiers are pseudonymized for processing.	Reduces avoidable differences in spelling and format while limiting the use of raw customer data.
2	KYC and relationship information is used to create a shortlist. Signals include occupation, seniority, corporate email domain, declared income or account vintage, metropolitan location, and existing HNI flags.	Avoids searching the full customer base and keeps external-data use focused.
3	Potential public profiles are searched using combinations such as full name plus employer, full name plus corporate domain, or email handle plus location.	A combination of details produces stronger candidates than a name-only search.
4	Each candidate is checked against employer, handle, location, career timeline, and name rarity. Strong matches continue automatically, uncertain matches go to an analyst, and weak matches are discarded.	Prevents an influential profile from being assigned to the wrong customer.
5	Confirmed profiles are assessed for reach, professional authority, engagement, and relevance. The resulting service tier and supporting reasons are sent to the CRM or ticketing workflow.	Turns the analysis into a clear service action instead of leaving it as a dashboard score.

Information stored in the service record

The operational record contains the customer identifier, match score, review status, influence score, service tier, supporting reasons, source timestamps, and any manual override. It does not need to contain a full copy of the external profile. This keeps the service team focused on the decision and reduces unnecessary exposure of personal information.

How customer matches and confidence scores are calculated

For every candidate profile, the workflow checks five types of evidence. Each check produces a value between 0 and 1: 1 means the detail agrees strongly, 0 means it does not agree, and a partial value is used where the match is close but not exact. The value is then multiplied by the importance assigned to that check. The five contributions are added to produce the final match score.

$$\text{Match score} = 0.35E + 0.20U + 0.20G + 0.15T + 0.10N$$

Code	Check	Maximum contribution	How it is checked
E	Employer or organisation	35 points	Compares the KYC employer with professional history, corporate information, or verified directorships.
U	Username or email handle	20 points	Compares the customer email handle with the public username or profile identifier.
G	Geographic location	20 points	Checks whether the city or metropolitan area agrees with the available customer information.
T	Career timeline	15 points	Checks whether age, education, employment, and career milestones form a plausible timeline.
N	Name rarity	10 points	Gives more weight to distinctive names and requires stronger supporting evidence for common names.

Worked example

Assume the employer matches exactly ($E = 1.00$), the email handle is a strong but not exact match ($U = 0.90$), the city matches ($G = 1.00$), the career timeline is mostly consistent ($T = 0.80$), and the name-rarity value is 0.60. The calculation is $0.35 + 0.18 + 0.20 + 0.12 + 0.06 = 0.91$. This candidate is treated as a confirmed match because the final score is above 0.85.

How the final score is used

Final score	Decision	What happens next
0.85 to 1.00	Confirmed match	The profile moves to influence assessment. The individual evidence values and source timestamps are retained with the decision.
0.70 to 0.84	Analyst review	An analyst checks the profile, employer, location, handle, and timeline before accepting or rejecting the match.
Below 0.70	Rejected match	The candidate is not used for influence scoring or service prioritisation.

How customer influence is assessed

Identity matching answers whether the profile belongs to the customer. Influence assessment answers a different question: if this customer raises a service issue publicly, how far could it travel and how relevant would that audience be to AxisDirect? Keeping these calculations separate prevents a weak identity match from receiving a high service tier simply because the candidate profile has a large following.

Composite Influence Score

$$\text{CIS} = \text{MIN}(100, [0.40 \times \text{Reach} + 0.35 \times \text{Authority} + 0.25 \times \text{Engagement}] \times \text{Relevance})$$

Component	Weight	What it measures
Reach	40%	The size of the customer's audience across the permitted sources. A logarithmic scale prevents extremely large accounts from dominating the score.
Authority	35%	Professional seniority, verified roles, company directorships, and recognised media or institutional standing.
Engagement	25%	Recent replies, shares, comments, and other interactions that show whether the audience actively responds.
Relevance	Multiplier	The relationship between the audience and financial markets. Finance and trading audiences receive a higher multiplier than unrelated audiences.

Customer groups identified in the project

Customer group	Typical indicators	Likely escalation route	Why the case may require attention
Startup founders and technology CXOs	Senior role, corporate domain, metropolitan base, large professional network	LinkedIn, X, or direct executive tagging	Issues can spread quickly through founder and technology networks.
Finance creators and trading educators	Active finance or trading content and a responsive retail audience	YouTube, Instagram, X, or private trading groups	Platform downtime or order problems can immediately affect audience confidence.
Financial journalists and editors	Editorial role at a business or financial publication	Editorial networks, public threads, or direct media inquiry	A customer complaint may develop into a wider story about service reliability.
Senior legal and regulatory professionals	Senior legal, regulatory, tax, or advisory position	Formal notices, board escalation, regulator, or ombudsman	The issue may require legal, compliance, or governance review.
Derivatives community leaders	Leader of a high-frequency trading group or private community	Group broadcasts and coordinated support escalation	One incident can generate many simultaneous tickets and complaints.

How service priority is applied

The influence score is converted into four service tiers so that support teams receive a clear instruction rather than another number to interpret. Each tier defines the queue, responsible team, response target, escalation point, and closure review. The tier does not change the customer's contractual rights or reduce the standard service available to other customers.

Tier	Representative customers	Service workflow	Response and resolution targets
Tier 1 - Executive VIP	Listed CXOs, nationally prominent finance creators, senior business editors, and major founders	Senior concierge desk; dedicated relationship owner; Head of CX notification; senior review before closure	First response under 15 minutes; resolution target under 4 hours; active monitoring during major incidents
Tier 2 - High priority	Growth-stage founders, corporate VPs or directors, analysts, and active trading coaches	Prioritized queue; senior Level-2 support; supervisor review of open cases	First response under 1 hour; resolution target under 12 hours; escalation if unresolved at 6 hours
Tier 3 - Monitored	Advisers, senior accountants, technology leads, and mid-tier commentators	Standard queue with an awareness flag; automatic elevation after repeated reopening or escalation	First response under 4 hours; standard 24-hour resolution target
Standard - General retail	Customers without a verified public-influence indicator	Standard IVR or chatbot triage and round-robin support	Existing institutional service levels

How the service team uses the tier

Once a tier is assigned, it is added to the customer's service record. It determines which support queue receives the case, who owns it, how quickly the first response is expected, and when an unresolved case must be escalated.

Changing a tier. If new information shows that the assigned tier is no longer appropriate, a relationship manager can request a change. An authorized supervisor approves the update and records why it was made.

Handling major incidents. During a platform outage or other widespread service event, the tier helps the incident team decide which customers require direct updates or senior follow-up. It does not reduce service for other customers.

Reviewing decisions. The operations team periodically samples cases to confirm that the customer match, assigned tier, and service response were correct. Any error is corrected in the service record.

Privacy and operational controls

The workflow uses customer information and external profile data, so the controls around purpose, source access, retention, security, and human review are part of the solution itself. These controls should be reflected in the data flow and system configuration, not added later as policy language.

Purpose limitation

The workflow is designed for service prioritisation and incident management. The score is not used for sales targeting, product eligibility, credit decisions, margin restrictions, employee screening, or general surveillance. Restricting the purpose also limits which data needs to be collected and who needs access to the result.

External-source authorization

For every external source, the project should maintain the provider name, fields used, access method, contractual or API entitlement, permitted purpose, refresh schedule, retention period, and deletion requirement. If LinkedIn data was used, the case study should name the exact approved access programme rather than referring generally to the LinkedIn API.

Data minimization and security

Raw KYC information remains in the client-controlled environment. Only the details needed for matching enter the processing workflow, preferably through pseudonymous identifiers. Access is role-based and logged. Rejected candidates and expired enrichment are deleted according to policy, while exports and bulk access are restricted.

Human review and correction

Analysts reviewing uncertain matches can see the evidence used in the calculation and record why they accepted or rejected the profile. Incorrect matches can be corrected, and service tiers can be appealed by authorized employees. Error rates and outcomes should be reviewed across name groups, languages, regions, and customer categories.

Key operational risks

Risk	Example	Control
Incorrect identity match	A common name is linked to the wrong public profile.	Several identity checks, analyst review, and a correction process.
Outdated profile information	The person's employer, location, or audience has changed.	Timestamp each source, define an expiry period, and recheck important changes.
Unauthorized data use	Publicly visible information is assumed to be freely reusable.	Maintain a source register and confirm the permitted purpose and access rights.
Unequal service outcomes	Customers without public visibility consistently receive slower service.	Protect standard service levels and review response times, errors, and overrides by segment.
Use outside the approved purpose	The score is reused for sales, credit, or account restrictions.	Limit output fields, restrict access, and prohibit reuse in policy and system design.

Project outcomes and client validation

ILLUSTRATIVE SCENARIO ONLY - The funnel counts and performance figures on this page are synthetic planning assumptions. They are not verified AxisDirect results and must be replaced or client-approved before external use.

The engagement established a repeatable operating model for identifying influential customers, validating the customer-to-profile match, and converting the assessment into a clear service action.

Confirmed engagement outputs

Output	Delivered scope
Customer population	Approximately 50 lakh KYC records across retail, HNI, and trading relationships
Identification workflow	Five stages covering normalization, shortlisting, candidate discovery, identity checks, and service-tier assignment
Assessment model	Five identity checks, followed by a separate assessment of reach, professional authority, engagement, and financial-market relevance
Service operating model	Four tiers linked to routing, ownership, response expectations, escalation, and review

Illustrative processing funnel

Stage	What the stage represents	Illustrative figure
Customer population	KYC population available for initial screening	5,000,000 records
Rule-based shortlist	Records meeting the KYC and relationship-screening rules	126,400 records
Candidate profiles assessed	Potential public or professional profiles evaluated for a match	24,860 profiles
Analyst review	Ambiguous matches checked by an authorized reviewer	4,720 profiles
Confirmed tiered profiles	High-confidence matches plus reviewed matches approved for the service workflow	8,140 profiles

Illustrative pilot results

Measure	Illustrative result
Confirmed-match accuracy	93.2% across a labelled review sample of 1,000 cases
Analyst-review rate	19.0% of candidate profiles, or 4,720 of 24,860 profiles assessed
Research time	Median identification time reduced from 24 minutes to 6 minutes
Service response	Median first-response time for priority cases improved by 38%

Client validation

"The workflow gave our service teams a consistent way to identify influential customers, understand why a case required additional attention, and route sensitive issues to the right team earlier."

Proposed wording only - client name, title, and written approval are required before publication.